

# Exercise 7: Connecting Astronomer Cloud with PostgreSQL

## Introduction

PostgreSQL is a popular open-source relational database, and connecting it to Airflow allows for automating database operations. This lab will guide you through configuring a connection to a PostgreSQL database hosted on [Render.com](https://render.com), a free managed service, and executing basic operations like creating and dropping tables using the provided Airflow DAG.

## Objectives

- Set up a free PostgreSQL database on Render.com.
- Configure a PostgreSQL connection in Astronomer Cloud.
- Deploy and run an Airflow DAG to create and drop a table in the PostgreSQL database.
- Understand the use of the `PostgresOperator` in Airflow.

## Preparation

### Prerequisites

- Familiarity with Airflow DAGs and task dependencies.
- Basic understanding of PostgreSQL and SQL commands.
- Docker installed and running.
- Python (3.7 or above) and pip installed.
- Active Astronomer Cloud account and workspace.
- Astro CLI installed on your local machine
  - Verify using bash command `astro version`

## Create a Free PostgreSQL Database on Render.com

- Sign up or log in to Render.com.
- Go to New + > PostgreSQL.
- Fill in the details:
- Name: Choose a name for your database.
- Region: Select a region closest to you.
- Free Plan: Ensure you select the free plan.
- Save the database and copy the connection details (host, port, database name, username, password).

## Let's Get Started

### Part 1: Environment Setup

#### Step 1. Authenticate Astro CLI:

- Authenticate your CLI with Astronomer Cloud:
  - i. Use command `astro login`
  - ii. Press Enter

#### Step 3. Project setup

- Ensure you have an initialized Airflow project. If not, initialize a new project using:
  - i. Create a new directory and navigate into it:
    1. `mkdir astro_cloud`
    2. `cd astro_cloud`
    3. `astro dev init`

### Part 2: Creating Postgres Connection in Astronomer Cloud

#### Step 1. Set Up PostgreSQL Connection in Astronomer Cloud

- Log in to your Astronomer Cloud workspace.
- Navigate to your Deployment and go to the Environment Variables or Connections tab.
- Get connection details from render.com
  - i. Go to dashboard
    1. Top right click on connect dropdown
    2. Select External
    3. Copy External Database URL available there in below format

```
postgresql://USER:PASSWORD@EXTERNAL_HOST:PORT/DATABASE
```

4. Use these parament later while creating code
- Add a new connection with the following details:
    - Connection ID: postgress\_render
    - Connection Type: Postgres
    - Host: EXTERNAL\_HOST from Render.com
    - Schema: DATABASE name from Render.com
    - Login: USER from Render.com
    - Password: PASSWORD from Render.com
    - Port: Default is 5432 unless specified otherwise by Render.

Save the connection.

## Install dbeaver on your machine

- Download and install dbeaver exe from <https://dbeaver.io/download/>
- Open dbeaver
- Go to File -> New > DBeaver -> Database connection ->PostgreSQL.
- Fill in the details: under Main
- Host: EXTERNAL\_HOST from Render.com
- Database: DATABASE name from Render.com
- Login: USER from Render.com
- Password: PASSWORD from Render.com
- Then Test connection and save

## Part 3: Deploying to Astronomer Cloud

### Step 1. Add a new Dag

- Navigate to the `dags/` directory and create a file named `Postgres_Dag.py`

```
from airflow import DAG

from airflow.providers.postgres.operators.postgres import PostgresOperator

from airflow.utils.dates import days_ago

from airflow.utils.dates import timedelta

# Define the default arguments for the DAG
default_args = {

    'owner': 'airflow',

    'retries': 1,

    'retry_delay': timedelta(minutes=5),

}

# Define the DAG
with DAG(

    'create_drop_postgres_table',

    default_args=default_args,

    description='A simple DAG to create and drop PostgreSQL table',

    schedule_interval=None, # Trigger manually or set cron schedule
```

```

        start_date=days_ago(1),

        catchup=False,

) as dag:

    # SQL query to create the table

    create_table_sql = """

    CREATE TABLE IF NOT EXISTS test_table (

        id SERIAL PRIMARY KEY,

        name VARCHAR(100),

        created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP

    );

    """

    # SQL query to drop the table

    drop_table_sql = "DROP TABLE IF EXISTS test_table;"

    # Task to create the table

    create_table = PostgresOperator(

        task_id='create_table',

        postgres_conn_id='postgres_render', # Connection ID set in Airflow UI

        sql=create_table_sql,

    )

    # Task to drop the table

```

```

drop_table = PostgresOperator(

    task_id='drop_table',

    postgres_conn_id='postgres_render', # Connection ID set in Airflow UI

    sql=drop_table_sql,

)

# Set the task dependencies

create_table >> drop_table

```

- Save the file and navigate to the project root.

## Step 2. Create deployment

- Use below command to create astor deployment
  - `astro deploy`
  - It will ask to create a deployment if no deployment is available as of now  
Provide details like
    - name of deployment
    - Cloud region where it need to be deployed
  - if deployment is available then select that deployment and proceed with  
deploy

## Step 3. Verify deployment

- Log in to the Astronomer Cloud UI and navigate to your deployment.
- Ensure your DAG appears under the DAGs tab.

#### Step 4. Monitor and Manage Workflows

- Trigger the `create_drop_postgres_table` DAG from the Astronomer Cloud UI.
- Monitor task execution:
  - **Task 1 (`create_table`)**: Creates the `test_table` in the PostgreSQL database.
  - **Task 2 (`drop_table`)**: Drops the `test_table`.

Check the Render.com database logs

Verify using SQL editor on dbeaver

#### Step 5. Update deployment

- Modify the DAG to include error for SQL execution.
- Log detailed error messages using Airflow's logging capabilities.
- Check failure logs on Airflow UI

## Conclusion

In this lab, you:

- Connected Astronomer Cloud to a PostgreSQL database hosted on Render.com.
- Configured and tested a DAG that performed basic database operations using the `PostgresOperator`.
- Explored key Airflow concepts like connections and task dependencies.

This exercise demonstrates how Airflow can be used to manage database operations programmatically, a critical skill for data engineers.